ATTENTION: CONTROL WHAT CAN INTERACT

Same frozen Qwen model. Different contexts. No fine-tuning.

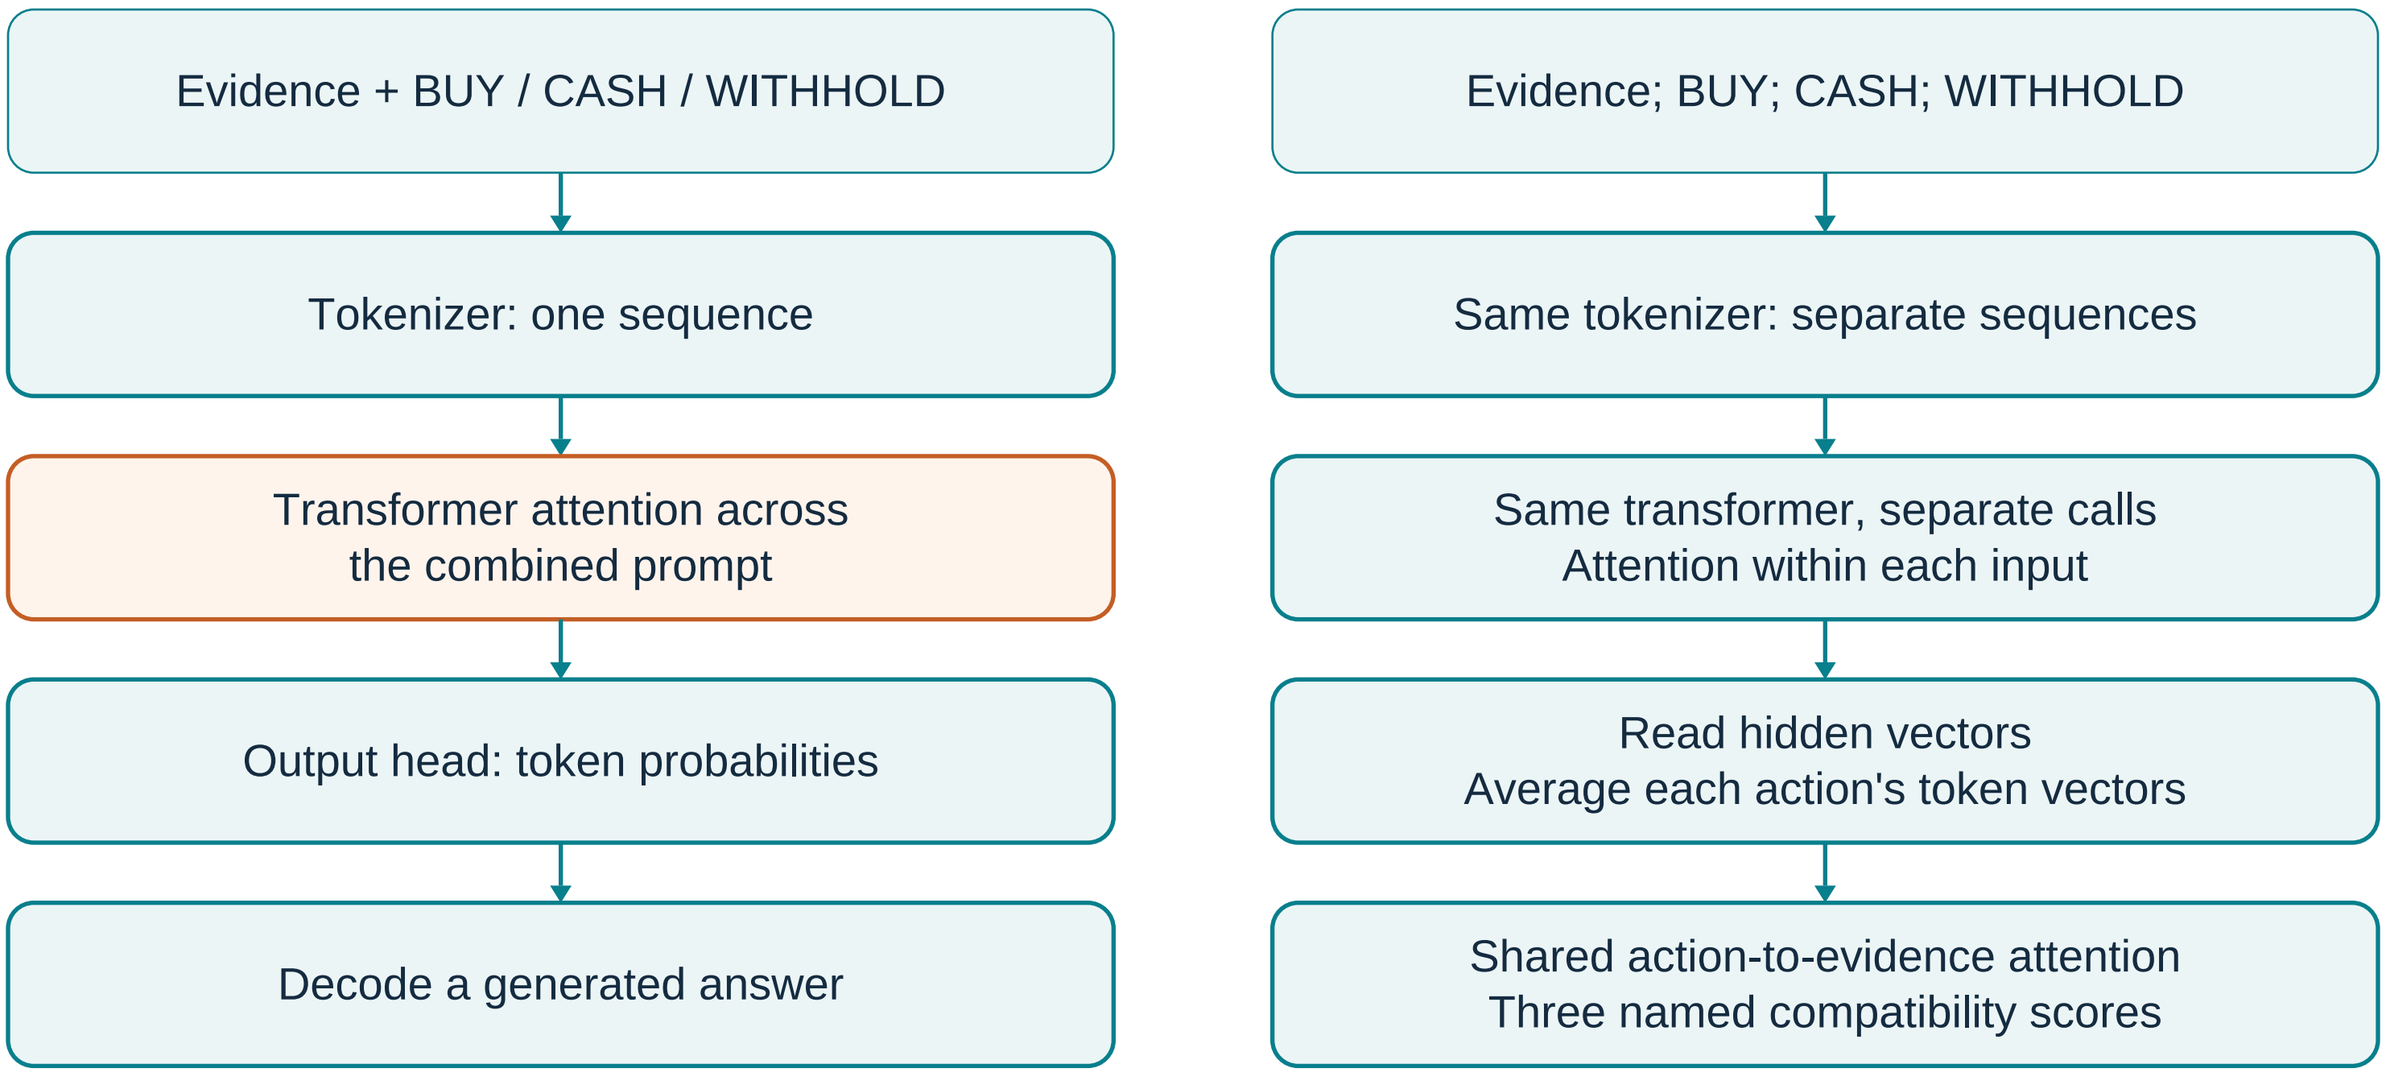

We bypass answer generation, not attention inside the transformer.

MEASURED: Qwen2.5-0.5B | 5 cases x 6 orders x 5 repeats

Order sensitivity	JSON in prompt	Separate scoring
Changed top-ranked action	32 / 75 pairs	0 / 75 pairs
Largest distribution shift (TV)	0.5975	0
Identical-repeat variation	0	0

Prompt results use scored answer probabilities, not generated-answer tests.
Separate scores are untrained compatibility values. Evidence order was not tested.

Matthew Dixon, PhD, FRM | BeliefLens

WORKING PAPER